

- A.) Raul Perez-Lopez (811921958)
- B.) The program has two mandatory inputs:
- The number of processes
 - The file being processed

Each unique process is represented using the **plist_t** struct provided

```
typedef struct {
    int pid;
    long offset;
    long lengthAssigned;
    int pipefd[2];
} plist_t;
```

I create a list of (number of processes) to store the processes information. Then, in the initial for loop I do the following:

- Compute the offset for each process
- Assign the length to that process

The lengthAssigned for each process is calculated by dividing the size of the file being processed by the number of processes we are using.

```
fseek(fp, 0L, SEEK_END);
fsize = ftell(fp);
fclose(fp);
long sizePerProcesss = fsize / numJobs; // Work load per child
```

Then, in the initial for loop I can assign each process their offset and assigned length.

```
for(i = 0; i < numJobs; i++) {
    plist[i].offset = sizePerProcesss * i;           // Calculate offset
    plist[i].lengthAssigned = sizePerProcesss;      // assign length to process
```

If the size of the file is NOT divisible by the number of processes, I assign the last process a little extra work to process the extra bytes.

```
// Assign the rest of the work (remainder) to last job if needed
if ((i == numJobs - 1) && remainder) {
    plist[i].lengthAssigned += remainder;
}
```

Where remainder here is the size of the file modulo with the number of processes.

I then set the pipe, using the **pipefd** variable like so:

```
// Set pipe
if (pipe(plist[i].pipefd) == -1) {
    perror("Error creating pipe.");
    exit(EXIT_FAILURE); // Exit with failure status
}
```

After assigning this information, I use `fork()` to launch a child process -> process its “workload” using the provided wordcount function -> if successful, save the results to the pipe (`fd[1]`).

```
} else if(pid == 0) { // Child
    fp = fopen(argv[2], "r");
    count = word_count(fp, plist[i].offset, plist[i].lengthAssigned);

    // send the result to the parent through pipe
    write(plist[i].pipefd[1], &count, sizeof(count_t)); // Write result to pipe
    close(plist[i].pipefd[1]); // Close write end of pipe
    fclose(fp);
    return 0;
}
```

This approach assigns equal units of “workload” to each individual process, with the last process having an extra ($N - 1$) extra units in the worst case.

For IPC, as previously stated I use a pipe. Where the child process writes to the write end of the pipe (`fd[1]`) and the parent reads from the read end of the pipe (`fd[0]`).

C.) For crash handling, we make sure that after launching the children processes that the parent waits for each individual process and ensures they terminated normally. I use a for loop, getting the PID for each child process from the `plist_t` list and checking the status using **WIFSIGNALED(status)**. While True, then we continue to launch a child process to make sure it correctly processes its “workload”.

```

// While the child process continues to exit irregularly
while (WIFSIGNALED(status)) {

    // Launch new process to replace this
    pid = fork();
    if (pid < 0) printf("Fork failed.\n");
    else if (pid == 0) {
        fp = fopen(argv[2], "r");
        count = word_count(fp, plist[i].offset, plist[i].lengthAssigned);

        write(plist[i].pipefd[1], &count, sizeof(count_t)); // Write result to pipe
        close(plist[i].pipefd[1]); // Close write end of pipe
        fclose(fp);
        return 0;
    }
    plist[i].pid = pid;
    waitpid(plist[i].pid, &status, 0);
}

```

Once the child terminated successfully, we can read the result from the pipe and add it to our total.

```

read(plist[i].pipefd[0], &result, sizeof(count_t)); // Parent reads result written by child
close(plist[i].pipefd[0]); // Close read end of pipe

// Add results to total
total.linecount += result.linecount;
total.wordcount += result.wordcount;
total.charcount += result.charcount;

```